

- When conducting the work of Lab 11, we conducted the test that uses the Central Limit Theorem even though the sample size was “small” (i.e., $n < 30$). It turns out, that how “far off” the t -test is can be computed using a first-order Edgeworth approximation for the error. Below, we will do this for the the further observations.

- Boos and Hughes-Oliver (2000) note that

$$P(T \leq t) \approx F_Z(t) + \underbrace{\frac{\text{skew}}{\sqrt{n}} \frac{(2t^2 + 1)}{6}}_{\text{error}} f_Z(t),$$

where $f_Z(\cdot)$ and $F_Z(\cdot)$ are the Gaussian PDF and CDF and skew is the skewness of the data. What is the potential error in the computation of the p -value when testing $H_0 : \mu_X = 0; H_a : \mu_X < 0$ using the zebra finch further data?

Solution: The potential error in the computation of the p -value is $-1.226006e - 13$

```
finch.dat = read.csv('./zebrafinches.csv')


#potential error in p-value
n = length(finch.dat$furthere)
t.furthere = t.test(finch.dat$furthere, alternative = 'less',
                    mu = 0)$statistic[1]
gauss.pdf = dnorm(t.furthere)
finch.skew = skewness(finch.dat$furthere)

error = gauss.pdf*finch.skew*(2*t.furthere^2 + 1)/(sqrt(n)*6)
```

- Compute the error for t statistics from -10 to 10 and plot a line that shows the error across t . Continue to use the skewness and the sample size for the zebra finch further data.

Solution:

```
#t-statistic error from -10 to 10
t.errors = c()
for (i in seq(-10, 10, by = 0.01)){
  t = i
  gauss.pdf = dnorm(t)
  error = gauss.pdf*finch.skew*(2*t^2 + 1)/(sqrt(n)*6)
  t.errors = append(t.errors, error)
}
par(mar = c(4, 4, 2, 1))
plot(seq(-10, 10, by = 0.01), t.errors,
     xlab = 't statistic',
     ylab = 'error in t',
     main = 'A changing error in t statistic')
```



- (c) Suppose we wanted to have a tail probability within 10% of the desired $\alpha = 0.05$. Recall we did a left-tailed test using the further data. How large of a sample size would we need? That is, we need to solve the error formula equal to 10% of the desired left-tail probability:

$$0.10\alpha \stackrel{\text{set}}{=} \underbrace{\frac{\text{skew}}{\sqrt{n}} \frac{(2t^2 + 1)}{6} f_Z(t)}_{\text{error}},$$

which yields

$$n = \left(\frac{\text{skew}}{6(0.10\alpha)} (2t^2 + 1) f_Z(t) \right)^2.$$

Solution: Calculating n with the given formula, we get a value very close to zero (on the order of 10^{-20}). That means that our data is so far from the null hypothesis of 0, that any datapoint over here will be beyond that 0.05 threshold.

```
#Sample size needed
a= 0.05
gauss.pdf = dnorm(t.further)
n = ((finch.skew/(6*0.10*a))*(2*t.further^2 + 1)*gauss.pdf)^2
```

2. Complete the following steps to revisit the analyses from lab 11 using the bootstrap procedure.

- (a) Now, consider the zebra finch data. We do not know the generating distributions for the closer, further, and difference data, so perform resampling to approximate the sampling distribution of the T statistic:

$$T = \frac{\bar{x}_r - 0}{s/\sqrt{n}},$$

where \bar{x}_r is the mean computed on the r^{th} resample and s is the sample standard deviation from the original samples. At the end, create an object called `resamples.null.closer`, for example, and store the resamples shifted to ensure they are consistent with the null hypotheses at the average (i.e., here ensure the shifted resamples are 0 on average, corresponding to $t = 0$, for each case).

Solution: In the below code, the mean of the shifted `resamples.null.closer` is approximately 0

```
R <- 10000
resamples <- tibble(xbars = rep(NA, R))

for(i in 1:R){
  curr.resample <- sample(finch.dat$'closer',
                        size = nrow(finch.dat),
                        replace = T)

  resamples$xbars[i] <- mean(curr.resample)
}

#Shifted Resamples
delta = mean(resamples$xbars)
resamples.null.closer = data.frame(xbars.shifted = resamples$xbars - delta)
```

- (b) Compute the bootstrap p -value for each test using the shifted resamples.

Solution: We calculated the bootstrap p -value with the below code. The output for the closer data is 1e-04, the further data is 1e-04, and the difference data is also 1e-04. This is probably due to an error in the code not working correctly, which I did not have time to rectify before the submission deadline.

```
####Further####
R <- 10000
resamples <- tibble(xbars = rep(NA, R))

for(i in 1:R){
  curr.resample <- sample(finch.dat$'further',
                        size = nrow(finch.dat),
                        replace = T)

  resamples$xbars[i] <- mean(curr.resample)
}

# Confidence Interval
quantile(resamples$xbars, c(0.025, 0.975))

library(boot)
boot.mean <- function(d, i){
  mean(d[i])
}
boots <- boot(data = finch.dat$'further',
             statistic = boot.mean,
             R = R)
boot.ci(boots, type="bca")
```

```

# shift so H0 is true
mean(resamples$xbars)
(delta <- mean(resamples$xbars) - 2.9)

resamples <- resamples |>
mutate(xbars.shifted = xbars - delta)

low <- 2.9 - delta
high <- 2.9 + delta

resamples |>
summarize(mean = mean(xbars.shifted),
           p.low = mean(xbars.shifted <= low),
           p.high = mean(xbars.shifted >= high))|>
mutate(p = p.low + p.high) |>
view()

boot.pval(boots, theta_null = mu0) # this does something slightly different

```

How do these compare to the t -test p -values?

- (c) What is the 5th percentile of the shifted resamples under the null hypothesis? Note this value approximates $t_{0.05, n-1}$. Compare these values in each case.

Solution: The outer 5th percentile of the shifted resamples are for the closer: 0.1218 and 0.1948, for further it is -0.2640 and -0.1606. Finally, for the difference, it is 0.2870 and 0.4509. We used the following code to find these values.

```

R <- 10000
resamples <- tibble(xbars = rep(NA, R))

for(i in 1:R){
  curr.resample <- sample(finch.dat$`diff`,
                        size = nrow(finch.dat),
                        replace = T)

  resamples$xbars[i] <- mean(curr.resample)
}

# Confidence Interval
quantile(resamples$xbars, c(0.025, 0.975))

library(boot)
boot.mean <- function(d, i){
  mean(d[i])
}
boots <- boot(data = finch.dat$`diff`,
             statistic = boot.mean,
             R = R)
boot.ci(boots, type="bca")

```

- (d) Compute the bootstrap confidence intervals using the resamples. How do these compare to the t -test confidence intervals?

Solution: The confidence interval is the same as these values for the outer 2.5% of the data.

3. Complete the following steps to revisit the analyses from lab 11 using the randomization procedure.

- (a) Now, consider the zebra finch data. We do not know the generating distributions for the closer, further, and difference data, so perform the randomization procedure

Solution: We did the following randomization procedure for each of the closer, further, and difference zebra finch data.

```

mu0 = 0
#Closer
R <- 10000
rand <- tibble(xbars = rep(NA, R))

# PREPROCESSING: shift the data to be mean 0 under H0
x.shift <- finch.dat$'closer' - mu0
# RANDOMIZE / SHUFFLE
for(i in 1:R){
  curr.rand <- x.shift *
    sample(x = c(-1, 1),
           size = length(x.shift),
           replace = T)

  rand$xbars[i] <- mean(curr.rand)
}
# Thinking is hard
rand <- rand |>
  mutate(xbars = xbars + mu0) # shifting back

```

- (b) Compute the randomization test p -value for each test.

Solution: To compute the randomization p -value for each text, we used the following code:

```

# p-value
(delta <- abs(mean(finch.dat$'closer') - mu0))
(low <- mu0 - delta) # mirror
(high <- mu0 + delta) # xbar

mean(rand$xbars <= low) +
mean(rand$xbars >= high)

```

The result is a p value that is effectively 0, which checks out with our other texts (ex, a plain t test) which have a p value on order of 10^{-8} .

- (c) Compute the randomization confidence interval by iterating over values of μ_0 .

Hint: You can “search” for the lower bound from Q_1 and subtracting by 0.0001, and the upper bound using Q_3 and increasing by 0.0001. You will continue until you find the first value for which the two-sided p -value is greater than or equal to 0.05.

Solution: For the closer data, the confidence interval is: 0.1162231 to 0.1962231. For the further data, it is -0.2627244 to -0.1427244. Finally, for the difference data, the confidence interval is 0.2689475 to 0.4489475. These intervals were found using the below code for the randomization test.

```

#Confidence Interval
mu0.iterate <- 0.01
starting.point <- mean(finch.dat$'closer')

mu.lower <- starting.point
repeat{
  rand <- tibble(xbars = rep(NA, R))

  # PREPROCESSING: shift the data to be mean 0 under H0
  x.shift <- finch.dat$'closer' - mu.lower
  # RANDOMIZE / SHUFFLE
  for(i in 1:R){
    curr.rand <- x.shift *
      sample(x = c(-1, 1),
             size = length(x.shift),
             replace = T)

    rand$xbars[i] <- mean(curr.rand)
  }
  # Thinking is hard
  rand <- rand |>
    mutate(xbars = xbars + mu.lower) # shifting back

  # p-value
  (delta <- abs(mean(finch.dat$'closer') - mu.lower))
  (low <- mu.lower - delta) # mirror
  (high <- mu.lower + delta) # xbar
  (p.val <- mean(rand$xbars <= low) +
    mean(rand$xbars >= high))
}

```

```

if(p.val < 0.05){
  break
}else{
  mu.lower <- mu.lower - mu0.iterate
}
}

mu.upper <- starting.point
repeat{
  rand <- tibble(xbars = rep(NA, R))

  # PREPROCESSING: shift the data to be mean 0 under H0
  x.shift <- finch.dat$'closer' - mu.upper
  # RANDOMIZE / SHUFFLE
  for(i in 1:R){
    curr.rand <- x.shift *
      sample(x = c(-1, 1),
            size = length(x.shift),
            replace = T)

    rand$xbars[i] <- mean(curr.rand)
  }
  # Thinking is hard
  rand <- rand |>
    mutate(xbars = xbars + mu.upper) # shifting back

  # p-value
  (delta <- abs(mean(finch.dat$'closer') - mu.upper))
  (low <- mu.upper - delta) # mirror
  (high <- mu.upper + delta) # xbar
  (p.val <- mean(rand$xbars <= low) +
    mean(rand$xbars >= high))

  if(p.val < 0.05){
    break
  }else{
    mu.upper <- mu.upper + mu0.iterate
  }
}

c(mu.lower, mu.upper)

```

4. **Optional Challenge:** In this lab, you performed resampling to approximate the sampling distribution of the T statistic using

$$T = \frac{\bar{x}_r - 0}{s/\sqrt{n}}.$$

I'm curious whether it is better/worse/similar if we computed the statistics using the sample standard deviation of the resamples (s_r), instead of the original sample (s)

$$T = \frac{\bar{x}_r - 0}{s_r/\sqrt{n}}.$$

- Perform a simulation study to evaluate the Type I error for conducting this hypothesis test both ways.
- Using the same test case(s) as part (a), compute bootstrap confidence intervals and assess their coverage – how often do we ‘capture’ the parameter of interest?

References

Boos, D. D. and Hughes-Oliver, J. M. (2000). How large does n have to be for z and t intervals? *The American Statistician*, 54(2):121–128.